

Nadeem Theba

Rajkot, India | +91 97149 65149 | nadeemtheba8@gmail.com | github.com/NADEEMTHEBA8 | linkedin.com/in/nadeem-theba

August 2026

Dear Hiring Manager,

I am writing to express my strong interest in a Data Engineer position at your organisation. With 3+ years of hands-on experience architecting scalable ETL/ELT pipelines, event-driven streaming architectures, and Medallion Lakehouses on AWS and Databricks, I am eager to contribute to data platform initiatives that drive measurable business impact.

At Kaival Technologies, I engineered daily ETL pipelines processing over 500GB of transactional data, reducing batch runtimes from 12 hours to 7.2 hours while cutting cloud compute overhead by 25%. I additionally established automated data validation frameworks using Pandera and Great Expectations that prevented 99.9% of anomalies from reaching production dashboards. These outcomes were achieved through careful architecture decisions in PySpark, Delta Lake, and Apache Airflow.

My project portfolio further demonstrates this expertise: I designed and built a Predictive Supply Chain Telemetry Pipeline ingesting real-time IoT sensor data across 50 machines and 5 global plants via AWS Kinesis, Databricks Auto Loader, and Delta Lake 3.x, reducing query scan volumes by 90% through Liquid Clustering. I also built an event-driven Real-Time Fraud Feature Store using Apache Kafka, PySpark Structured Streaming, Redis, and FastAPI that delivers feature vectors in sub-10ms latency.

I bring strong proficiency in Python, PySpark, SQL, dbt, Kafka, DuckDB, Terraform, and the AWS ecosystem (S3, Kinesis, Lambda, RDS), and I am experienced in enforcing Shift-Left data quality through automated schema contracts and CI/CD pipelines via GitHub Actions and Docker. I welcome the opportunity to discuss how my experience can contribute to your engineering team.

Sincerely,

Nadeem Theba

Data Engineer | nadeemtheba8@gmail.com