Nadeem Theba

Rajkot, India | +91 97149 65149 | nadeemtheba8@gmail.com | github.com/NADEEMTHEBA8 |
linkedin.com/in/nadeem-theba-602862208

PROFESSIONAL SUMMARY

Data Engineer with 3+ years of experience architecting scalable ETL/ELT pipelines, event-driven streaming architectures, and Medallion Lakehouses on AWS & Databricks. Highly proficient in **Databricks**, **PySpark**, **SQL**, **Delta Lake**, **dbt**, **AWS (Lambda, CloudWatch, S3, Kinesis)**, and **Terraform.** Experienced in ingesting high-frequency IoT sensor telemetry, constructing real-time customer routing logic, and optimizing physical lakehouse storage layouts.

TECHNICAL SKILLS

- **Programming & Querying:** Python (pandas, scikit-learn, Pydantic, GenAI APIs), SQL, Spark SQL
- **Data Engineering & Lakehouse:** Databricks, PySpark, Delta Lake 3.x, Medallion Architecture, dbt Core, DuckDB, ETL/ELT
- **Streaming & Processing:** Apache Kafka, Amazon Kinesis, Debezium CDC, PySpark Structured Streaming, Redis
- **Cloud & Infrastructure:** AWS (S3, Kinesis, Lambda, CloudWatch, RDS, IAM, PostgreSQL, Snowflake), Docker, Terraform, CI/CD

PROFESSIONAL EXPERIENCE

Data Engineer | Kaival Technologies | *Remote (London / India) | 10/2023 – Present*

- Engineered scalable ETL pipelines utilizing Python and SQL to process over 500GB of daily data, successfully reducing complex batch runtimes from 12hrs to 7.2hrs and cutting compute costs by 25%.
- Implemented strict data quality rules using programmatic validation, successfully preventing 99.9% of anomalies from reaching production dashboards.
- Architected scalable data integration workflows using Python, consolidating 250k+ legacy records from disparate enterprise source systems into a centralized schema.
- Achieved zero-downtime production releases by architecting automated CI/CD pipelines via GitHub Actions & Docker.

TECHNICAL EXPERIENCE

supply-chain-telemetry-pipeline | *AWS S3, Kinesis, Databricks, PySpark, Delta Lake, dbt, Terraform | 01/2024 – 04/2024*

- Ingested real-time IoT sensor telemetry across 50 machines and 5 global plants into an S3 Medallion Lakehouse, calculating 24h rolling risk scores to prevent unplanned downtime.
- Eliminated silent streaming data loss by implementing deep response inspection and exponential backoff with jitter on Kinesis batch writes for throughput exceptions.
- Reduced analytical query scan volumes by 90% via physical Delta Lake partitioning on plant location and event date.

realtime-fraud-feature-store | *Apache Kafka, PySpark Streaming, Redis, FastAPI, dbt, Docker | 09/2023 – 12/2023*

- Architected an event-driven streaming feature store using Kafka and PySpark Structured Streaming to parse raw payloads into Delta Lake with exact-once checkpointing.
- Served complete historical feature vectors to a FastAPI customer routing and decisioning endpoint in sub-10ms latency by decoupling PostgreSQL storage with Redis.
- Ensured 100% pipeline resilience by building automated payload routing to a Kafka Dead-Letter Queue (DLQ) for replay.

coverdrive | *DuckDB, dbt, PySpark, Pandera, AWS S3, Apache Airflow | 05/2023 – 08/2023*

- Mitigated executor memory skew during high-cardinality joins by engineering custom PySpark key-salting algorithms across distributed partitions.
- Prevented downstream pipeline corruption by 99.9% through Shift-Left Pandera data contracts enforcing non-null constraints in Airflow.
- Optimized lakehouse analytics queries by running dbt transformations directly on DuckDB querying Parquet on S3 at sub-second execution cost.

EDUCATION

Master's Degree, Data Science and Analytics | University of Hertfordshire, UK | *09/2021 – 09/2023*

Bachelor of Technology (B.Tech), Computer Science | Marwadi University, India | *08/2016 – 05/2020*